MST & Shortest Path

Now that we all know what a graph is, its basic properties, and how to traverse them. Now let's talk about some more specific features of graphs.

Minimum Spanning Tree (MST)

A tree is a perennial plant with an elongated stem, or trunk, supporting... Oh wait, wrong wikipedia page.

A tree in CS is a **connected**, **acyclic** graph. A spanning tree in a graph is sub-graph that is a tree that connects all vertices within a given graph. A **Minimum Spanning Tree (MST)** is a spanning tree with the minimum total edge weight.

In English, a MST is a subset of the edges that connects the graph with the minimum total weight. There are various real life applications of minimum spanning tree, such as choosing how to lay optic cable to connect a bunch of data centers or roads connecting cities. But in contest, a lot of times you will just be asked to "find the MST."

There are two algorithms used, both are pretty intuitive, and have the same runtime but make sure you understand the essence because sometimes you'll have to construct the MST without constructing the graph.

Prim's Algorithm

Prim's algorithm starts with an arbitrary vertex, v_0, and divide up the graph into two sets of vertices: V and U. Initially $V = \{v_0\}$ and U is everything else in the graph. While U is non-empty / $V \neq G$, we pick the edge (v, u) with $v \in V$ and $u \in U$ with the minimum edge weight. Add the edge to the MST, and move u form U to V. In psudocode, it is:

```
Prim(Graph G):
    V <- v_0                            // arbitrary starting point
    U = G \ V
    MST = {}                            // initialize to empty set
    while V != G.V:                     // while not all vertices are in V
        e = (u, v) with minimum edge weight where u in U and v in V
        MST.add(e)
        U.remove(u)
        V.add(u)
    return MST
```

The runtime of this largely depends on finding the edge with minimum weight between the visited and unvisited set. If we conduct a linear search, the runtime is $O(V^2)$ whereas if we use a priority queue it is $O(E \log(V))$.

Kruskal's Algorithm

Kruskal's algorithm is the opposite of Prim's. Instead of building a tree spawning from one vertex. It builds a forest that eventually becomes one single spanning tree. This algorithm is greedy in nature as well. We began by picking the shortest edge, adding that to the MST, then we combine the two trees at the endpoints of the edge.

Shortest Path – Dijkstra

As I have said last time, there is a reason why DFS and BFS should be renamed to depth first transversal and breath first transversal because there are better searching algorithms. And today we're going to learn about the better algorithm – Dijkstra's. The advantage of this over the other two is that this algorithm is targeted in finding the shortest path between two points in a graph, and have significantly better run time. Dijkstra's algorithm is a modified version of what is known as **best first search**.

We assign a distance to each node, 0 at the initial node and infinity at every other node. Then at each node, we calculate the tentative distance to each of its neighbors by adding the weight of the edge to the distance at the current node. If the tentative distance is smaller than the distance currently associated with the neighbor, we will update the distance to the smaller value. We will then mark the current node as visited and move to the neighbor with the lowest associated distance.

The algorithm terminates when the node we are looking for has been visited, or when the entire graph has been traversed. The algorithm has the worst-case runtime of $O(E + V \log(V))$ if we use a priority queue to store distance.

Floyd-Warshall

Dijkstra's algorithm is one that calculates the shortest distance from one point to another (and potentially one point to every other point in the graph). However, sometimes we need to find all pairwise distances between all pairs of nodes. This can be done with repeated Dijkstra's with time complexity of $O(EV + V^2 \log(V))$. However, this is doing quite a bit of overcounting as we are traversing the graph more times than necessary. In truth we only need to traverse all the edges one per vertex. This is known as the Floyd-Warshall algorithm. It repeated gets estimates of all pairwise distances and stop at the optimal one.

In psudocode it is:

```
let dist be a 2D array of distances between vertices initialized to infinity
for v from 1...V:
    dist[v][v] = 0
for (u, v) in E with weight w:
    dist[u][v] = w
```

```
for i from 1...V:
    for j from 1...V:
        for k from 1...V:
            if dist[i][j] > dist[i][k] + dist[k][j]
                dist[i][j] = dist[i][k] + dist[k][j]
```

Basically the idea behind this is that if we can go from point i to j faster through k, we'll do it. Note that this algorithm does not give the exact path, a simple modification will return the path. [every time we do the substitution w/ k, we know it is a part of the path]